

Apache Hive

Working with Apache Hive

During this exercise, you become an Apache Hive Developer/Analyst who will create physical objects in Hive and perform some basic operations on them.

Select The Username

Please use your Linux account at the edge node (cdh00.cl.ii.pw.edu.pl)

Using Hive command line interface

1 Login using ssh to edgenode cdh00.cl.ii.pw.edu.pl using your Linux account.

2. In this exercise we are going to use the CSV file we uploaded to HDFS earlier today. First of all check if the file is in your HDFS home folder:

```
hdfs dfs -ls
/user/${USER}/external/measured_data/measured_data.csv
```

3a. Before you can access secured (by Kerberos) distributed file system you have to generate Kerberos ticket. When prompted provide your password and afterwards verify that ticket has been generated.

```
kinit
```

Password for <USERNAME>@CL.II.PW.EDU.PL:

```
klist
```

3b Login to hive and create your own Hive database

```
hive
hive> create database ${USER};
```

4. Open another ssh connection or if you are using terminal manager like byobu/tmux/screen just open another window. Check if your database exists:

```
hive -e "show databases;" 2>/dev/null | grep ${USER}
```

In this step you executed a hive command directly from Linux cmd line. What do you think – where can it be useful?

5. Create an external Hive table in your database using a CSV file you have in your HDFS home folder:

```
CREATE EXTERNAL TABLE MEASURED_DATA_CSV_EXTERNAL
(
  `md_lsb_id` int,
  `md_timestamp` string,
  `md_lsb_id2` int,
  `md_value` double,
  `md_unit` string,
  `md_timetype` string,
  `md_quality_mark` string,
  `md_desc` string
)
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' LINES TERMINATED BY '\n' STORED as TEXTFILE
LOCATION '/user/${USER}/external/measured_data';
```

6. Change your working database to yours and list all the tables “stored” in it:

```
use ${USER};
OK
Time taken: 0.318 seconds

hive> show tables;
OK
measured_data_csv_external
```

7. Print and analyze table metadata:

```
hive> desc formatted measured_data_csv_external;
```

```
OK
```

# col_name	data_type	comment
md_lsb_id	int	
md_timestamp	string	
md_lsb_id2	int	
md_value	double	
md_unit	string	
md_timetype	string	
md_quality_mark	string	
md_desc	string	

```
# Detailed Table Information
```

```
Database:          xmwiewio
Owner:             xmwiewio
CreateTime:        Thu May 26 13:33:52 CEST 2016
LastAccessTime:    UNKNOWN
Protect Mode:      None
Retention:         0
Location:
hdfs:///cdh01.cl.ii.pw.edu.pl:802/user/xmwiewio/external/measured_data
Table Type:        EXTERNAL_TABLE
Table Parameters:
    COLUMN_STATS_ACCURATE    false
    EXTERNAL                  TRUE
    numFiles                  1
    numRows                   -1
    rawDataSize               -1
    totalSize                 331035256
    transient_lastDdlTime     1464262432
```

```
# Storage Information
```

```
SerDe Library:
org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe
InputFormat:      org.apache.hadoop.mapred.TextInputFormat
OutputFormat:
org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat
Compressed:       No
Num Buckets:      -1
Bucket Columns:   []
Sort Columns:     []
Storage Desc Params:
    field.delim      |
```

```
line.delim          \n
serialization.format |
Time taken: 0.389 seconds, Fetched: 41 row(s)
```

8. Generate a DDL (data definition) for a table:

```
hive>show create table measured_data_csv_external;
OK
CREATE EXTERNAL TABLE `training.measured_data_csv_external` (
  `md_lsb_id` int,
  `md_timestamp` string,
  `md_lsb_id2` int,
  `md_value` double,
  `md_unit` string,
  `md_timetype` string,
  `md_quality_mark` string,
  `md_desc` string)
ROW FORMAT DELIMITED
  FIELDS TERMINATED BY '|'
  LINES TERMINATED BY '\n'
STORED AS INPUTFORMAT
  'org.apache.hadoop.mapred.TextInputFormat'
OUTPUTFORMAT
  'org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat'
LOCATION
  'hdfs://HDFS-HA/user/sar_wim/external/measured_data'
TBLPROPERTIES (
  'COLUMN_STATS_ACCURATE'='false',
  'numFiles'='1',
  'numRows'='-1',
  'rawDataSize'='-1',
  'totalSize'='331035256',
  'transient_lastDdlTime'='1464262432')
```

```
Time taken: 0.276 seconds, Fetched: 25 row(s)
```

9. Try to run a simple query to get an average value of md_value column grouped by md_lsb_id and a year and month for the last few months;

```
hive> SELECT md_lsb_id,date_format(md_timestamp,'yyyy-MM') as
year_month,avg(md_value) as avg_val FROM
measured_data_csv_external GROUP BY
md_lsb_id,date_format(md_timestamp,'yyyy-MM') ORDER BY
md_lsb_id,year_month desc limit 15;
```

OK

```
1      2017-05 0.49873148036001064
1      2016-05 0.49768154708603957
1      2015-05 0.497515021264509
2      2017-05 0.4988676805140444
2      2016-05 0.49742044523741197
2      2015-05 0.5056076317793062
3      2017-05 0.49836022891096265
3      2016-05 0.4991462674179214
3      2015-05 0.500562383337185
4      2017-05 0.4993031206736664
4      2016-05 0.4955264848762207
4      2015-05 0.49834381017309315
5      2017-05 0.4983890152540982
5      2016-05 0.498885413611568
5      2015-05 0.4989340026906018
Time taken: 69.092 seconds, Fetched: 15 row(s)
```

10. Create a copy of the measured_data_csv_external table as managed by Hive and stored in ORC format. In another ssh session or terminal window compare the sizes of both tables:

```
hive>CREATE TABLE MEASURED_DATA_ORC STORED AS ORC AS SELECT * FROM
MEASURED_DATA_CSV_EXTERNAL;
```

```
hdfs dfs -du -h -s /user/${USER}/external/measured_data
482.0 M  1.4 G  /user/xmwiewio/external/measured_data
```

```
hdfs dfs -du -h -s
/warehouse/tablespace/managed/hive/${USER}.db/measured_data_or
c
52.4 M  157.3 M
/warehouse/tablespace/managed/hive/xmwiewio.db/measured_data_o
```

rc

11. Rerun the query from step 9 using ORC table and compare the time taken:

```
hive> SELECT md_lsb_id,date_format(md_timestamp,'yyyy-MM') as  
year_month,avg(md_value) as avg_val FROM measured_data_orc GROUP  
BY md_lsb_id,date_format(md_timestamp,'yyyy-MM') ORDER BY  
md_lsb_id,year_month desc limit 15;
```

OK

```
1      2017-05 0.4987314803600108  
1      2016-05 0.4976815470860367  
1      2015-05 0.497515021264509  
2      2017-05 0.49886768051404695  
2      2016-05 0.4974204452374088  
2      2015-05 0.5056076317793062  
3      2017-05 0.4983602289109608  
3      2016-05 0.499146267417921  
3      2015-05 0.500562383337185  
4      2017-05 0.49930312067366467  
4      2016-05 0.4955264848762195  
4      2015-05 0.49834381017309315  
5      2017-05 0.4983890152540996  
5      2016-05 0.498885413611569  
5      2015-05 0.4989340026906018  
Time taken: 52.808 seconds, Fetched: 15 row(s)
```

What do you think – what are the reasons for the speedup, if any?

12. Create a table partitioned by year:

```

CREATE TABLE MEASURED_DATA_ORC_PART
(
  `md_lsb_id` int,
  `md_timestamp` string,
  `md_lsb_id2` int,
  `md_value` double,
  `md_unit` string,
  `md_timetype` string,
  `md_quality_mark` string,
  `md_desc` string
)
PARTITIONED BY (dt int)
STORED AS ORC;

```

```
hive> desc MEASURED_DATA_ORC_PART;
```

```
OK
```

```

md_lsb_id          int
md_timestamp        string
md_lsb_id2          int
md_value            double
md_unit             string
md_timetype         string
md_quality_mark     string
md_desc            string
dt                  int

```

```
# Partition Information
```

```
# col_name          data_type          comment
```

```
dt                  int
```

```
Time taken: 0.411 seconds, Fetched: 14 row(s)
```

13. Load the year partitions using Hive “INSERT OVERWRITE PARTITION”

```

INSERT OVERWRITE TABLE MEASURED_DATA_ORC_PART
PARTITION(dt=2015)
SELECT * FROM MEASURED_DATA_ORC WHERE year(md_timestamp)=2015;

```

```

INSERT OVERWRITE TABLE MEASURED_DATA_ORC_PART
PARTITION(dt=2016)
SELECT * FROM MEASURED_DATA_ORC WHERE year(md_timestamp)=2016;

```

```

INSERT OVERWRITE TABLE MEASURED_DATA_ORC_PART
PARTITION(dt=2017)

```

```
SELECT * FROM MEASURED_DATA_ORC WHERE year(md_timestamp)=2017;
```

14. List table partitions:

```
hive> show partitions measured_data_orc_part;
OK
dt=2013
dt=2014
dt=2015
```

15. Overwrite all the partitions at once using dynamic partitioning option:

```
set hive.exec.dynamic.partition.mode=nonstrict;
INSERT OVERWRITE TABLE MEASURED_DATA_ORC_PART PARTITION(dt)
SELECT md_lsb_id,
md_timestamp,
md_lsb_id2,
md_value,
md_unit,
md_timetype,
md_quality_mark,
md_desc,
year(md_timestamp) as dt FROM MEASURED_DATA_ORC;
```

```
Time taken for load dynamic partitions : 505
```

```
Loading partition {dt=2017}
```

```
Loading partition {dt=2016}
```

```
Loading partition {dt=2015}
```

```
Time taken for adding to write entity : 1
```

```
Partition xmviewio.measured_data_orc_part{dt=2015} stats:
[numFiles=1, numRows=2400000, totalSize=18330755,
rawDataSize=1140000000]
```

```
Partition xmviewio.measured_data_orc_part{dt=2016} stats:
[numFiles=1, numRows=2400000, totalSize=18331378,
rawDataSize=1140000000]
```

```
Partition xmviewio.measured_data_orc_part{dt=2017} stats:
[numFiles=1, numRows=2400000, totalSize=18330712,
rawDataSize=1140000000]
```

```
MapReduce Jobs Launched:
```

```
OK
```


16. Compare the execution time, CPU time and HDFS reads of non-partitioned and partitioned table:

```
hive> select count(*) from measured_data_orc_part where dt=2015
and md_lsb_id=100;
```

```
Hadoop job information for Stage-1: number of mappers: 1;
number of reducers: 1
2017-05-18 21:39:51,763 Stage-1 map = 0%, reduce = 0%
2017-05-18 21:39:57,951 Stage-1 map = 100%, reduce = 0%,
Cumulative CPU 1.36 sec
2017-05-18 21:40:04,128 Stage-1 map = 100%, reduce = 100%,
Cumulative CPU 3.19 sec
MapReduce Total cumulative CPU time: 3 seconds 190 msec
Ended Job = job_1495125662310_0016
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.19 sec
HDFS Read: 57176 HDFS Write: 6 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 190 msec
OK
24000
Time taken: 20.56 seconds, Fetched: 1 row(s)
-----
hive> select count(*) from measured_data_orc where
year(md_timestamp)=2015 and md_lsb_id=100;
```

```
Hadoop job information for Stage-1: number of mappers: 1;
number of reducers: 1
2017-05-18 21:40:16,918 Stage-1 map = 0%, reduce = 0%
2017-05-18 21:40:23,099 Stage-1 map = 100%, reduce = 0%,
Cumulative CPU 1.96 sec
2017-05-18 21:40:29,291 Stage-1 map = 100%, reduce = 100%,
Cumulative CPU 3.99 sec
MapReduce Total cumulative CPU time: 3 seconds 990 msec
Ended Job = job_1495125662310_0017
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.99 sec
HDFS Read: 169556 HDFS Write: 6 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 990 msec
OK
```

24000

Time taken: 19.509 seconds, Fetched: 1 row(s)